

Elastisearch

Create a 3 node elasticsearch cluster with the name pp_elasticsearch.

STEPS:-

1. Add the hostnames in all the 3 nodes in **/etc/hosts** path.
2. Install the following dependencies using the below commands (to install elasticsearch.)

- apt update
- apt-get install default-jre

```
sre@stg-janvimadhu003:~$ sudo su
Verification code:
root@stg-janvimadhu003:/home/sre# apt update
|Hit:1 http://mirror.phonepe.nb6/focal focal InRelease
|Hit:2 http://mirror.phonepe.nb6/focal focal-updates InRelease
|Hit:3 http://mirror.phonepe.nb6/focal focal-backports InRelease
|Hit:4 http://mirror.phonepe.nb6/focal focal-security InRelease
|Get:5 http://mirror.phonepe.nb6/pkgs InRelease [3278 B]
|Hit:6 http://mirror.phonepe.nb6/saltstack/latest focal InRelease
|Hit:7 http://repo.powerdns.com/ubuntu focal-rec-44 InRelease
|Ign:5 http://mirror.phonepe.nb6/pkgs InRelease
Fetched 3278 B in 1s (4397 B/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
All packages are up to date.
W: GPG error: http://mirror.phonepe.nb6/pkgs InRelease: The following signatures couldn't be verified because the public key is not available: NO_PUBKEY 3B95A5FF6824D053A
root@stg-janvimadhu003:/home/sre# apt-get install default-jre
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  at-spi2-core ca-certificates-java default-jre-headless fontconfig-config
  fonts-dejavu-core fonts-dejavu-extra java-common libatk-bridge2.0-0
  libatk-wrapper-java libatk-wrapper-java-jni libatk1.0-0 libatk1.0-data
  libatspi2.0-0 libavahi-client3 libavahi-common-data libavahi-common3
  libbcpus2 libdrm-amdgpu1 libdrm-intel1 libdrm-nouveau2 libdrm-radeon1
  libfontconfig1 libfontenc1 libgif7 libgl1 libgl1-mesa-dri libglapi-mesa
  libglvnd0 libglx-mesa0 libglx0 libgraphite2-3 libharfbuzz0b libice6
  libjpeg-turbo8 libjpeg8 liblcms2-2 libllvm12 libpcieaccess0 libpcsc-lite1
  libssensors-config libssensors5 libsm libvulkan1 libwayland-client0
  libxi1-xcb1 libxaw7 libxcb-dr12-0 libxcb-dr13-0 libxcb-glx0 libxcb-present0
  libxcb-randr0 libxcb-shape0 libxcb-shm0 libxcb-sync1 libxcb-xfixes0
  libcomposite1 libfixes3 libft2 libxi libxinerama1 libxxkbfile1 libxmu6
  libxpm4 libxrandr2 libxrender1 libxshmfence1 libxt6 libxtst6 libxv1
  libxf86dg1 libxf86vm1 mesa-vulkan-drivers openjdk-11-jre
  openjdk-11-jre-headless x11-common x11-utils
```

- sudo apt-get install apt-transport-https
- wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add -
- echo "deb <https://artifacts.elastic.co/packages/6.x/apt> stable main" | sudo tee -a /etc/apt/sources.list.d/elastic-6.x.list

- **sudo apt-get update**

```
root@stg-janvimadhu003:/home/sre# wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add -
OK
root@stg-janvimadhu003:/home/sre# echo "deb https://artifacts.elastic.co/packages/6.x/apt stable main" | sudo tee -a /etc/apt/sources.list.d/elastic-6.x.list
deb https://artifacts.elastic.co/packages/6.x/apt stable main
root@stg-janvimadhu003:/home/sre# sudo apt-get update
Hit:1 http://mirror.phonepe.nb6/focal focal InRelease
Hit:2 http://mirror.phonepe.nb6/focal focal-updates InRelease
Hit:3 http://mirror.phonepe.nb6/focal focal-backports InRelease
Hit:4 http://mirror.phonepe.nb6/focal focal-security InRelease
Get:5 http://mirror.phonepe.nb6/pkgs InRelease [3278 B]
Hit:6 http://mirror.phonepe.nb6/salstack/latest focal InRelease
Hit:7 http://repo.powerdns.com/ubuntu focal-rec-44 InRelease
Get:8 https://artifacts.elastic.co/packages/6.x/apt stable InRelease [7123 B]
Ign:5 http://mirror.phonepe.nb6/pkgs InRelease
Get:9 https://artifacts.elastic.co/packages/6.x/apt stable/main amd64 Packages [85.4 kB]
Fetched 95.8 kB in 2s (49.2 kB/s)
Reading package lists... Done
W: GPG error: http://mirror.phonepe.nb6/pkgs InRelease: The following signatures couldn't be verified because the public key is not available: NO_PUBKEY 3B95A5FF6824D53A
root@stg-janvimadhu003:/home/sre# sudo apt-get install elasticsearch
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  elasticsearch
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 149 MB of archives.
After this operation, 240 MB of additional disk space will be used.
Get:1 https://artifacts.elastic.co/packages/6.x/apt stable/main amd64 elasticsearch all 6.8.23 [149 MB]
Fetching 149 MB in 2s (74.3 kB/s)
Selecting previously unselected package elasticsearch.
(Reading database ... 98705 files and directories currently installed.)
Preparing to unpack .../elasticsearch_6.8.23_all.deb ...
Creating elasticsearch group... OK
Creating elasticsearch user... OK
Unpacking elasticsearch (6.8.23) ...
Setting up elasticsearch (6.8.23) ...
Created elasticsearch keystore in /etc/elasticsearch
Processing triggers for systemd (245.4-4ubuntu3.11) ...
```

3. Install elastic search using the below command:

- **sudo apt install elasticsearch**

4. Do the following configuration on node1 in **/etc/elasticsearch/elasticsearch.yml** using the below command:

```

GNU nano 4.8                               /etc/elasticsearch/elasticsearch.yml
#
# ----- Cluster -----
#
# Use a descriptive name for your cluster:
#
cluster.name: pp_elasticsearch
#
# ----- Node -----
#
# Use a descriptive name for the node:
#
node.name: stg-janvimadhu002
node.master: true
node.data: true
#
# Add custom attributes to the node:
#
#node.attr.rack: r1
#
# ----- Paths -----
#
# Path to directory where to store the data (separate multiple locations by comma):
#
path.data: /var/lib/elasticsearch
#
# Path to log files:
#
path.logs: /var/log/elasticsearch
#
# ----- Memory -----
#
# Lock the memory on startup:
#
#bootstrap.memory_lock: true
#
# Make sure that the heap size is set to about half the memory available
# on the system and that the owner of the process is allowed to use this
# limit.
#
# Elasticsearch performs poorly when the system is swapping the memory.
#
# ----- Network -----
#
# Set the bind address to a specific IP (IPv4 or IPv6):
#
network.host: 10.57.55.41
#
# Set a custom port for HTTP:
#
http.port: 9200
#
# For more information, consult the network module documentation.
#
# ----- Discovery -----
#
# Pass an initial list of hosts to perform discovery when new node is started:
# The default list of hosts is ["127.0.0.1", "[::1]"]
#
discovery.zen.ping.unicast.hosts: ["10.57.55.41", "10.57.55.42", "10.57.55.43"]
#
# Prevent the "split brain" by configuring the majority of nodes (total number of master-eligible nodes / 2 + 1):
#
discovery.zen.minimum_master_nodes: 2
#

```

5. Do the same configuration on node2 and node3 with their respective network. host.

6. Restart **elasticsearch.service** to save the configuration on all the 3 nodes.

```
[sre@stg-janvimadhu002:~$ sudo systemctl start elasticsearch.service^C
[sre@stg-janvimadhu002:~$ sudo systemctl status elasticsearch.service
● elasticsearch.service - Elasticsearch
   Loaded: loaded (/lib/systemd/system/elasticsearch.service; disabled; v>
     Active: active (running) since Wed 2022-03-23 18:05:11 IST; 3min 37s a>
       Docs: http://www.elastic.co
      Main PID: 200123 (java)
        Tasks: 43 (limit: 4682)
       Memory: 1.2G
      CGroup: /system.slice/elasticsearch.service
              └─200123 /bin/java -Xms1g -Xmx1g -XX:+UseConcMarkSweepGC -XX:C>
                  └─200220 /usr/share/elasticsearch/modules/x-pack-ml/platform/l>

Mar 23 18:05:11 stg-janvimadhu002 systemd[1]: Started Elasticsearch.
Mar 23 18:05:11 stg-janvimadhu002 elasticsearch[200123]: OpenJDK 64-Bit Ser>
lines 1-13/13 (END)
```

7. Use the below command to check the cluster health (general info on the cluster and gauge its health)

```
root@stg-janvimadhu002:/home/sre# sudo curl http://10.57.55.41:9200
{
  "name" : "stg-janvimadhu002",
  "cluster_name" : "pp_elasticsearch",
  "cluster_uuid" : "GM1x672TQrqw3HJdB0XgVA",
  "version" : {
    "number" : "6.8.23",
    "build_flavor" : "default",
    "build_type" : "deb",
    "build_hash" : "4f67856",
    "build_date" : "2022-01-06T21:30:50.087716Z",
    "build_snapshot" : false,
    "lucene_version" : "7.7.3",
    "minimum_wire_compatibility_version" : "5.6.0",
    "minimum_index_compatibility_version" : "5.0.0"
  },
  "tagline" : "You Know, for Search"
}
root@stg-janvimadhu002:/home/sre# sudo curl 10.57.55.41:9200/_cluster/health?pretty
{
  "cluster_name" : "pp_elasticsearch",
  "status" : "green",
  "timed_out" : false,
  "number_of_nodes" : 3,
  "number_of_data_nodes" : 3,
  "active_primary_shards" : 0,
  "active_shards" : 0,
  "relocating_shards" : 0,
  "initializing_shards" : 0,
  "unassigned_shards" : 0,
  "delayed_unassigned_shards" : 0,
  "number_of_pending_tasks" : 0,
  "number_of_in_flight_fetch" : 0,
  "task_max_waiting_in_queue_millis" : 0,
  "active_shards_percent_as_number" : 100.0
}
```

1. Create an index with pp_index and note down the number of primary/replica shards created by default.

STEPS:

1. To create a index with pp_index we will use the below command:

```
sudo curl -XPUT "10.57.55.41:9200/pp_index?pretty"
```

2. To check the number of primary and replica shards created by default use the below command:

```
root@stg-janvimadhu002:/home/sre# sudo curl http://10.57.55.41:9200/_cat/shards
pp_index 3 p STARTED 0 261b 10.57.55.42 stg-janvimadhu003
pp_index 3 r STARTED 0 261b 10.57.55.43 stg-janvimadhu004
pp_index 4 r STARTED 0 261b 10.57.55.42 stg-janvimadhu003
pp_index 4 p STARTED 0 261b 10.57.55.43 stg-janvimadhu004
pp_index 2 r STARTED 0 261b 10.57.55.42 stg-janvimadhu003
pp_index 2 p STARTED 0 261b 10.57.55.41 stg-janvimadhu002
pp_index 1 p STARTED 0 261b 10.57.55.43 stg-janvimadhu004
pp_index 1 r STARTED 0 261b 10.57.55.41 stg-janvimadhu002
pp_index 0 p STARTED 0 261b 10.57.55.42 stg-janvimadhu003
pp_index 0 r STARTED 0 261b 10.57.55.41 stg-janvimadhu002
```

3. To check whether or not our index is created or not we can use the command given below:

```
root@stg-janvimadhu002:/home/sre# sudo curl -XGET "10.57.55.41:9200/pp_index?pretty"
{
  "pp_index" : {
    "aliases" : { },
    "mappings" : { },
    "settings" : {
      "index" : {
        "creation_date" : "1648041117362",
        "number_of_shards" : "5",
        "number_of_replicas" : "1",
        "uuid" : "PGUxXXUwTTmST8XAAtelSQ",
        "version" : {
          "created" : "6082399"
        },
        "provided_name" : "pp_index"
      }
    }
  }
}
```

2. Insert some data into the index by creating 20 to 30 documents.

Using the bulk api we can directly insert multiple data into our index using the below command:

```

root@stg-janvimadhu002:/home/sre# sudo curl -X POST "http://10.57.55.41:9200/_bulk" -H 'Content-Type: application/json' -d '{"index":{"_id": "4"} } {"name":"janvi"} {"index":{"_id": "5"} } {"name":"janvi"} > {"index":{"_id": "6"} } {"name":"janvi"} > {"index":{"_id": "7"} } {"name":"janvi"} > {"index":{"_id": "8"} } {"name":"janvi"} > {"index":{"_id": "9"} } {"name":"janvi"} > {"index":{"_id": "10"} } {"name":"janvi"} > {"index":{"_id": "11"} } {"name":"janvi"} > {"index":{"_id": "12"} } {"name":"janvi"} {"index":{"_id": "13"} } {"name": "janvi"} {"index":{"_id": "14"} } {"name": "janvi"} {"index":{"_id": "15"} } {"name": "janvi"} {"index":{"_id": "16"} } {"name": "janvi"} {"index":{"_id": "17"} } {"name": "janvi"} {"index":{"_id": "18"} } {"name": "janvi"} {"index":{"_id": "19"} } {"name": "janvi"} > {"index":{"_id": "20"} } {"name": "janvi"} >
{
  "took":36,"errors":false,"items":[{"_index":"pp_index","_type":"names","_id":4,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":0,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":5,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":1,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":6,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":2,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":7,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":3,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":8,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":4,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":9,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":5,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":10,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":6,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":11,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":7,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":12,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":8,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":13,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":9,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":14,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":10,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":15,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":11,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":16,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":12,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":17,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":13,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":18,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":14,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":19,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":15,"_primary_term":1,"status":201}, {"_index":"pp_index","_type":"names","_id":20,"_version":1,"result":"created","_shards":{"total":2,"successful":2,"failed":0},"_seq_no":16,"_primary_term":1,"status":201}
}

```

3. Read the data from the indice and write the same to a json file. [file name : read_data.json]

To read the data from an index we can use the below command:

```

root@stg-janvimadhu002:/home/sre# sudo curl -XGET "http://10.57.55.41:9200/_search?pretty"
{
  "took" : 18,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : 20,
    "max_score" : 1.0,
    "hits" : [
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "14",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "19",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "5",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "8",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "9",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "10",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "11",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "12",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "13",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "14",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "15",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "16",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "17",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "18",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "19",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      },
      {
        "_index" : "pp_index",
        "_type" : "names",
        "_id" : "20",
        "_score" : 1.0,
        "_source" : {
          "name" : "janvi"
        }
      }
    ]
  }
}

```

To write the same response onto a json file we can use elasticdump command:

```
root@stg-janvimadhu002:/home/sre# sudo elasticdump --input=http://10.57.55.41:9200/pp_index --output=read_file.json --type=data
[Wed, 23 Mar 2022 18:35:10 GMT | starting dump
Wed, 23 Mar 2022 18:35:10 GMT | got 20 objects from source elasticsearch (offset: 0)
Wed, 23 Mar 2022 18:35:10 GMT | sent 20 objects to destination file, wrote 20
[Wed, 23 Mar 2022 18:35:10 GMT | got 0 objects from source elasticsearch (offset: 20)
Wed, 23 Mar 2022 18:35:10 GMT | Total Writes: 20
Wed, 23 Mar 2022 18:35:10 GMT | dump complete]
```

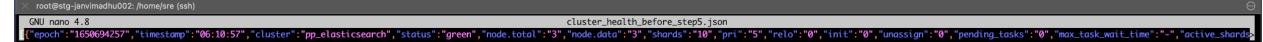
```
GNU nano 4.8                                read_file.json
[{"_index": "pp_index", "_type": "names", "_id": "14", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "19", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "5", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "8", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "9", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "10", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "12", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "2", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "4", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "6", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "15", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "20", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "1", "_score": 1, "_source": {"name": "Janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "7", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "13", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "16", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "18", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "3", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "11", "_score": 1, "_source": {"name": "janvi"}}, {"_index": "pp_index", "_type": "names", "_id": "17", "_score": 1, "_source": {"name": "janvi"}}]
```

4. Capture the elasticsearch cluster health along with a timestamp and write it to file. [file name : cluster_health_before_step5.json]

Now lets change our nodes configuration in such a way that stg-janvimadhu002 is master node and stg-janvimadhu003 and stg-prabhjotbedi004 are data nodes.

Now to capture ElasticSearch cluster health along with timestamp onto a file we can use the below command:

```
root@stg-janvimadhu002:/home/sre# sudo curl http://10.57.55.41:9200/_cat/health?format=json > cluster_health_before_step5.json
% Total    % Received % Xferd  Average Speed   Time     Time      Current
                                         Dload  Upload Total Spent   Left Speed
100    265  100    265    0     0  33125      0  --:--:-- --:--:-- --:--:-- 33125
```



```
root@stg-janvimadhu002:/home/sre (ssh)
GNU nano 4.8                                     cluster_health_before_step5.json
[{"epoch": "1650694257", "timestamp": "06:10:57", "cluster": "pp_elasticsearch", "status": "green", "node.total": "3", "node.data": "3", "shards": "10", "pri": "5", "relo": "0", "init": "0", "unassign": "0", "pending_tasks": "0", "max_task_wait_time": "", "active_shards": "3"}]
```

5. Stop two data nodes from the cluster and ensure you are still able to read the same data from es and write it to a json file. [file name : read_data_after_step5.json]

- Stop 2 data nodes using systemctl stop elasticsearch.service

```
root@stg-janvimadhu003:/home/sre# sudo systemctl stop elasticsearch.service
root@stg-janvimadhu003:/home/sre# sudo systemctl status elasticsearch.service
● elasticsearch.service - Elasticsearch
    Loaded: loaded (/lib/systemd/system/elasticsearch.service; disabled; vendor>
      Active: inactive (dead)
        Docs: http://www.elastic.co

Apr 23 10:57:33 stg-janvimadhu003 systemd[1]: Started Elasticsearch.
Apr 23 10:57:33 stg-janvimadhu003 elasticsearch[574383]: OpenJDK 64-Bit Server>
Apr 23 11:01:17 stg-janvimadhu003 systemd[1]: Stopping Elasticsearch...
Apr 23 11:01:17 stg-janvimadhu003 systemd[1]: elasticsearch.service: Succeeded>
Apr 23 11:01:17 stg-janvimadhu003 systemd[1]: Stopped Elasticsearch.
Apr 23 11:01:17 stg-janvimadhu003 systemd[1]: Started Elasticsearch.
Apr 23 11:01:18 stg-janvimadhu003 elasticsearch[574528]: OpenJDK 64-Bit Server>
Apr 23 11:07:47 stg-janvimadhu003 systemd[1]: Stopping Elasticsearch...
Apr 23 11:07:47 stg-janvimadhu003 systemd[1]: elasticsearch.service: Succeeded>
Apr 23 11:07:47 stg-janvimadhu003 systemd[1]: Stopped Elasticsearch.
lines 1-15/15 (END)
```

- Now when we will try to read data from es it will give a message all shards failed we can write this message in the form of json using the below command.(meaning we cannot read data if our data nodes fail)

```
root@stg-janvimadhu002:/home/sre# sudo curl "http://10.57.55.41:9200/_pp_index/_search?pretty" > read_data_after_step5.json
% Total    % Received % Xferd  Average Speed   Time     Time      Time  Current
          Dload  Upload Total Spent   Left Speed
100  2011  100  2011    0     0 14062      0 --:--:-- --:--:-- --:--:-- 14161
```

GNU nano 4.8								read_data_after_step5.json	
<pre>{ "error" : { "root_cause" : [], "type" : "search_phase_execution_exception", "reason" : "all shards failed", "phase" : "query", "grouped" : true, "failed_shards" : [] }, "status" : 503 } }</pre>									

6. Capture the elasticsearch cluster health along with a timestamp and write it to file. [file name : cluster_health_after_step5.json]

root@stg-janvimadhu002:/home/sre# sudo curl http://10.57.55.41:9200/_cat/health?format=json > cluster_health_after_step5.json									
<pre>% Total % Received % Xferd Average Speed Time Time Time Current Dload Upload Total Spent Left Speed 100 261 100 261 0 0 9321 0 --:--:-- --:--:-- --:--:-- 9666</pre>									
<pre>cluster_health_after_step5.json [{"epoch":1650694792,"timestamp":"06:19:52","cluster":"pp_elasticsearch","status":"red","node.total":1,"node.data":0,"shards":0,"pri":0,"relo":0,"ver":0}]"</pre>									

7. Add the nodes back to the cluster.

Use systemctl start elasticsearch.service command on the 2 data nodes to bring them back to the cluster

```
root@stg-janvimadhu003:/home/sre# systemctl start elasticsearch
root@stg-janvimadhu003:/home/sre# systemctl status elasticsearch
● elasticsearch.service - Elasticsearch
   Loaded: loaded (/lib/systemd/system/elasticsearch.service; disabled; vendor preset: enabled)
   Active: active (running) since Sat 2022-04-23 11:50:45 IST; 4s ago
     Docs: http://www.elastic.co
 Main PID: 580366 (java)
   Tasks: 19 (limit: 4682)
  Memory: 1.1G
    CGroup: /system.slice/elasticsearch.service
            └─580366 /bin/java -Xms1g -Xmx1g -XX:+UseConcMarkSweepGC -XX:CMSInitiatingOccupancyFraction=75 -XX:+UseCMSInitiatingOccupancyOnly -Des.networkaddress.cache.negative.ttl=0 -Djava.awt.headless=true -Djava.net.preferIPv4Stack=true -jar /usr/share/elasticsearch/modules/x-pack-ml/platform/linux-x86_64/bin/controller
```

8. Increase replica shard count to 3 (for each primary shard you need to have 3 replicas) and capture the location of primary and replica shards

1. To increase the replica shard count by 3 use the below command:

```
root@stg-janvimadhu002:/home/sre# sudo curl -X PUT "10.57.55.41:9200/pp_index/_settings?pretty" -H 'Content-Type: application/json' -d'
> {
>   "index" : {
>     "number_of_replicas" : 3
>   }
> }
{
  "acknowledged" : true
}
```

```
root@stg-janvimadhu002:/home/sre# sudo curl -X PUT "10.57.55.41:9200/pp_index/_settings?pretty" -H 'Content-Type: application/json' -d'
> {
>   "index" : {
>     "number_of_replicas" : 3
>   }
> }
{
  "acknowledged" : true
}
```

2. To capture the location of primary and replica shards use the below command (Here I have also made **stg-janvimadhu002** a data node so that maximum number of replica shards can be assigned, unassigned replica shards can get assigned when more data nodes are added to the cluster)

```
5}root@stg-janvimadhu002:/home/sre# sudo curl http://10.57.55.41:9200/_cat/shards
pp_index 3 r STARTED    5 6.7kb 10.57.55.43 stg-janvimadhu004
pp_index 3 p STARTED    5 6.7kb 10.57.55.42 stg-janvimadhu003
pp_index 3 r UNASSIGNED
pp_index 3 r UNASSIGNED
pp_index 1 p STARTED    5 3.5kb 10.57.55.43 stg-janvimadhu004
pp_index 1 r STARTED    5 3.5kb 10.57.55.42 stg-janvimadhu003
pp_index 1 r UNASSIGNED
pp_index 1 r UNASSIGNED
pp_index 4 r STARTED    3 6.7kb 10.57.55.43 stg-janvimadhu004
pp_index 4 p STARTED    3 6.7kb 10.57.55.42 stg-janvimadhu003
pp_index 4 r UNASSIGNED
pp_index 4 r UNASSIGNED
pp_index 2 r STARTED    5 6.7kb 10.57.55.43 stg-janvimadhu004
pp_index 2 p STARTED    5 6.7kb 10.57.55.42 stg-janvimadhu003
pp_index 2 r UNASSIGNED
pp_index 2 r UNASSIGNED
pp_index 0 r STARTED    2 3.5kb 10.57.55.43 stg-janvimadhu004
pp_index 0 p STARTED    2 3.5kb 10.57.55.42 stg-janvimadhu003
pp_index 0 r UNASSIGNED
pp_index 0 r UNASSIGNED
```